

project_code

July 20, 2020

0.1 # Image classification with Machine Learning

0.2 University of Milan

0.2.1 DataScience and Economics - Machine Learning Module

Authors : Andrea Ierardi, Emanuele Morales, Gregorio Luigi Saporito

How to load the dataset:

```
from google.colab import drive
drive.mount('/content/gdrive')

import os
os.environ['KAGGLE_CONFIG_DIR'] = "/content/gdrive/My Drive/Kaggle"

%cd /content/gdrive/My Drive/Kaggle

!kaggle datasets download --force -d moltean/fruits

!unzip fruits.zip
```

1 Image classification with Neural Networks

1.1 1. The dataset

1.1.1 1.1 Libraries

```
[ ]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import os
from tqdm import tqdm
import random
import pandas as pd

from plotnine import *
from sklearn.decomposition import PCA
```

```

from sklearn.datasets import load_files
from keras.preprocessing.image import array_to_img, img_to_array, load_img
from sklearn import preprocessing

from keras.utils import np_utils
from sklearn.utils import shuffle
import numpy as np

import matplotlib.pyplot as plt

from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D
from keras.layers import Activation, Dense, Flatten, Dropout
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import ModelCheckpoint
from keras import backend as K

from keras.applications import MobileNetV2

```

Using TensorFlow backend.

2 1.1 Data Loading

```

[ ]: DATADIR = "fruits-360/Training"
DATADIR_test = "fruits-360/Test"

TYPES = ["Apple", "Banana", "Plum", "Pepper", "Cherry", "Grape", "Tomato", "↵
↵"Potato", "Pear", "Peach"]
fruits = {}
def load_dataset(dire):
    fruits = {}
    images_as_array = []
    labels = []
    for category in tqdm(os.listdir(dire)):
        for typ in TYPES:
            if(category.split()[0] == typ):
                fruits[category] = typ
                path = os.path.join(dire, category)
                class_num = TYPES.index(fruits[category])

                class_name = fruits[category]
                for img in tqdm(os.listdir(path)):
                    file = os.path.join(path, img)

```

```

        images_as_array.
    ↪ append(img_to_array(load_img(file,target_size=(32, 32))))
        labels.append(class_num)
    images_as_array = np.array(images_as_array)
    labels = np.array(labels)
    return images_as_array, labels

```

2.0.1 Split in test and training sets

```
[ ]: train = load_dataset(DATADIR)
test = load_dataset(DATADIR_test)
```

```
[ ]: x_train, y_train= train
```

```
[ ]: x_test, y_test = test
```

2.0.2 Train and test shape

```
[ ]: print('Train shape:')
print('X: ',x_train.shape)
print('y: ',y_train.shape)

print('Test shape')
print('X: ',x_test.shape)
print('y: ',y_test.shape)
```

Train shape:

X: (32607, 32, 32, 3)

y: (32607,)

Test shape

X: (10906, 32, 32, 3)

y: (10906,)

3 1.2 Pre-processing

3.0.1 Pre-process the labels and the images

```
[ ]: x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255

no_of_classes = len(np.unique(y_train))
y_train = np_utils.to_categorical(y_train,no_of_classes)
y_test = np_utils.to_categorical(y_test,no_of_classes)

[ ]: print(y_train[0:10])
print("Number of classes: ",no_of_classes)
```

```

[[1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]

```


Number of classes: 10

3.0.2 Visualisation of the first 10 images

```

[ ]: fig = plt.figure(figsize =(30,5))
    for i in range(10):
        ax = fig.add_subplot(2,5,i+1,xticks=[],yticks=[])
        ax.imshow(np.squeeze(x_train[i]))
        ax.set_title("{} ".format(TYPES[np.
↪argmax(y_train[i]))],color="green"),fontdict= {'fontsize': '25'})

```



3.0.3 Suffle of the data

```

[ ]: x_train,y_train = shuffle(x_train, y_train)
    x_test,y_test = shuffle(x_test, y_test)


```

3.0.4 Visualisation of the first 10 images shuffled

```

[ ]: fig = plt.figure(figsize =(30,5))
    for i in range(10):
        ax = fig.add_subplot(2,5,i+1,xticks=[],yticks=[])
        ax.imshow(np.squeeze(x_train[i]))
        ax.set_title("{} ".format(TYPES[np.
↪argmax(y_train[i]))],color="green"),fontdict= {'fontsize': '25'})

```



3.0.5 Split in validation and test set

```
[ ]: # Using 80-20 rule
split = len(x_test)*80//100

print('Test len before split: ',len(x_test))
print('Validation split len:', split)
```

Test len before split: 10906
Validation split len: 8724

```
[ ]: # Now, we have to divide the validation set into test and validation set
x_test,x_valid = x_test[split:],x_test[:split]
y_test,y_valid = y_test[split:],y_test[:split]
print('Train X : ',x_train.shape)
print('Train y : ',y_train.shape)

print('1st training image shape ',x_train[0].shape)

print('Validation X : ',x_valid.shape)
print('Validation y : ',y_valid.shape)
print('Test X : ',x_test.shape)
print('Test y : ',y_test.shape)
```

Train X : (32607, 32, 32, 3)
Train y : (32607, 10)
1st training image shape (32, 32, 3)
Validation X : (8724, 32, 32, 3)
Validation y : (8724, 10)
Test X : (2182, 32, 32, 3)
Test y : (2182, 10)

3.0.6 Definition of zero-one loss function

```
[ ]: def zero_one(prediz,test):
    y_hat = []
    y_t = []
    for i in range(len(prediz)):
```

```

y_hat.append(np.argmax(prediz[i]))
y_t.append(np.argmax(test[i]))

loss = []
for i in range(len(prediz)):
    if(y_hat[i] == y_t[i]):
        loss.append(0)
    else:
        loss.append(1)

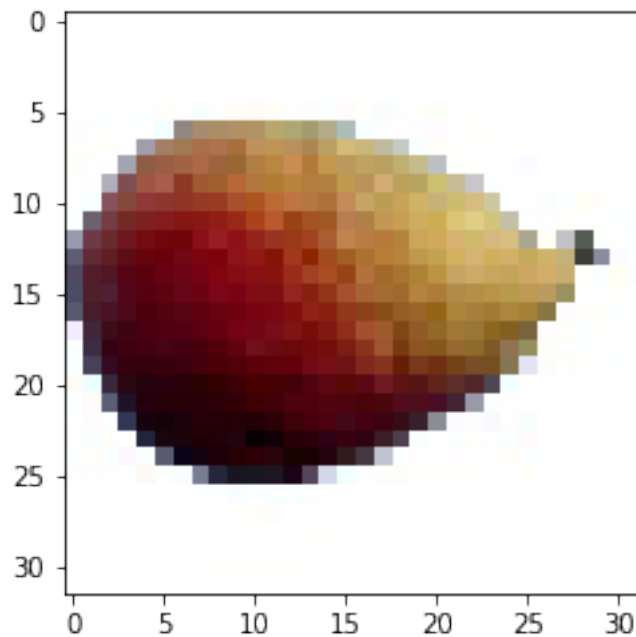
return np.mean(loss)

```

4 1.3 PCA and feed-forward NN

```
[ ]: plt.imshow(x_train[0])
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f72992944a8>
```




```

[ ]: x_train[0].shape
type(x_train[1])
rgb_weights = [0.2989, 0.5870, 0.1140]
image_test = x_train[0]
image_grey = np.dot(image_test[..., :3], rgb_weights)

```

```
plt.imshow(image_grey, cmap=plt.get_cmap("gray"))
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f729920a198>
```



```
[ ]: # transform np.ndarray from rgb to grey
x_train_grey = np.ndarray(shape=(x_train.shape[0], 32, 32))
for i in range(x_train.shape[0]):
    image_convert = x_train[i]
    x_train_grey[i] = np.dot(image_convert[:, :, 3], rgb_weights)

x_valid_grey = np.ndarray(shape=(x_valid.shape[0], 32, 32))
for i in range(x_valid.shape[0]):
    image_convert = x_valid[i]
    x_valid_grey[i] = np.dot(image_convert[:, :, 3], rgb_weights)

x_test_grey = np.ndarray(shape=(x_test.shape[0], 32, 32))
for i in range(x_test.shape[0]):
    image_convert = x_test[i]
    x_test_grey[i] = np.dot(image_convert[:, :, 3], rgb_weights)
```

```
[ ]: # flatten 32x32 images by concatenating them into a vector, each column of the
      ↪ matrix will be an image
x_train_flat = np.ndarray(shape=(1024, x_train_grey.shape[0]))
for i in range(x_train_grey.shape[0]):
    x_train_flat[:, i] = np.concatenate(x_train_grey[i])
```

```

x_valid_flat = np.ndarray(shape=(1024, x_valid_grey.shape[0]))
for i in range(x_valid_grey.shape[0]):
    x_valid_flat[:,i] = np.concatenate(x_valid_grey[i])

x_test_flat = np.ndarray(shape=(1024, x_test_grey.shape[0]))
for i in range(x_test_grey.shape[0]):
    x_test_flat[:,i] = np.concatenate(x_test_grey[i])

```

```

[ ]: standard_scaler = preprocessing.StandardScaler()
x_train_flat_T = standard_scaler.fit_transform(x_train_flat.T)
x_valid_flat_T = standard_scaler.transform(x_valid_flat.T)
x_test_flat_T = standard_scaler.transform(x_test_flat.T)

```

```

[ ]: x_train_flat_T.shape

```

```

[ ]: (32607, 1024)

```

```

[ ]: x_train_flat = x_train_flat_T.T
x_valid_flat = x_valid_flat_T.T
x_test_flat = x_test_flat_T.T

```

```

[ ]: x_train_flat.shape

```

```

[ ]: (1024, 32607)

```

```

[ ]: a = np.cov(x_train_flat)
b = np.linalg.eig(a)
b[0].shape

```

```

[ ]: (1024,)

```

```

[ ]: b

```

```

[ ]: (array([384.41591328, 86.57192039, 60.39849467, ..., 0.
, 0.
, 0.
]),
array([[ -2.77148730e-03,  1.99671963e-02, -7.00245270e-03, ...,
0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
[ 1.06770166e-04,  5.26867240e-03, -3.60037909e-03, ...,
0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
[ -1.31497722e-04,  4.57673250e-03, -2.40477865e-03, ...,
0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
...,
[ 9.24994377e-04,  3.61739014e-03,  2.95348978e-03, ...,
0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
[ 1.13335613e-03,  5.13559532e-03,  3.96259018e-03, ...,
0.00000000e+00,  0.00000000e+00,  0.00000000e+00],

```




```
[ 0.00000000e+00,  0.00000000e+00,  0.00000000e+00, ...,
 0.00000000e+00,  0.00000000e+00,  1.00000000e+00]])
```


4.0.1 PCA explained variance ratio and “Eigenfruits”

```
[ ]: pca = PCA().fit(x_train_flat)
plt.figure(figsize=(18, 7))
plt.plot(pca.explained_variance_ratio_.cumsum(), lw=3)
```

```
[ ]: [<matplotlib.lines.Line2D at 0x7f72991f9b38>]
```



```
[ ]: # try to plot some of the eigenvectors, the so called "eigenfruits"
fig = plt.figure(figsize=(30,5))
for i in range(10):
    ax = fig.add_subplot(2,5,i+1,xticks=[],yticks=[])
    ax.imshow(np.squeeze(b[1][:,i].reshape(32,32)))
```



4.0.2 Reduce image noise with PCA

```
[ ]: x_train_flat.shape, x_valid_flat.shape, x_test_flat.shape
```

```
[ ]: ((1024, 32607), (1024, 8724), (1024, 2182))
```

```
[ ]: def PCA_iter(x_all,start, end, step):
    lis = []
    for i in range(start, end, step):
        print("\n\n==== Component: ",i,"====\n")

        (train,valid, test) = x_all
        pca = PCA(n_components=i)
        print("original shape:  ", train.shape)

        pca.fit_transform(train)

        train_PCA = pca.transform(train)
        train_new = pca.inverse_transform(train_PCA)

        valid_PCA = pca.transform(valid)
        valid_new = pca.inverse_transform(valid_PCA)

        test_PCA = pca.transform(test)
        test_new = pca.inverse_transform(test_PCA)

        print("transformed shape:", train_PCA.shape)
        print("final shape:", train_new.shape)

        tupla = (x_train_PCA, x_valid_PCA, x_test_PCA)
        ↪=train_new,valid_new,test_new

        lis.append(tupla)

    return lis

[ ]: lis_PCA = PCA_iter((x_train_flat_T,x_valid_flat_T, x_test_flat_T),10,211,20)
```

```
==== Component:  10 ====
```

```
original shape:  (32607, 1024)
transformed shape: (32607, 10)
final shape: (32607, 1024)
```

```
==== Component:  30 ====
```

```
original shape:  (32607, 1024)
transformed shape: (32607, 30)
```

final shape: (32607, 1024)

===== Component: 50 =====

original shape: (32607, 1024)
transformed shape: (32607, 50)
final shape: (32607, 1024)

===== Component: 70 =====

original shape: (32607, 1024)
transformed shape: (32607, 70)
final shape: (32607, 1024)

===== Component: 90 =====

original shape: (32607, 1024)
transformed shape: (32607, 90)
final shape: (32607, 1024)

===== Component: 110 =====

original shape: (32607, 1024)
transformed shape: (32607, 110)
final shape: (32607, 1024)

===== Component: 130 =====

original shape: (32607, 1024)
transformed shape: (32607, 130)
final shape: (32607, 1024)

===== Component: 150 =====

original shape: (32607, 1024)
transformed shape: (32607, 150)
final shape: (32607, 1024)

===== Component: 170 =====

original shape: (32607, 1024)

```
transformed shape: (32607, 170)
final shape: (32607, 1024)
```

```
===== Component: 190 =====
```


```
original shape: (32607, 1024)
transformed shape: (32607, 190)
final shape: (32607, 1024)
```

```
===== Component: 210 =====
```

```
original shape: (32607, 1024)
transformed shape: (32607, 210)
final shape: (32607, 1024)
```


```
[ ]: # 10 components example of the same image
tr,va,te = lis_PCA[1]
plt.imshow(tr[2,:].reshape(32,32), cmap=plt.get_cmap("gray"))
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f7296f20a20>
```



```
[ ]: # 210 components example of an image
tr,va,te = lis_PCA[len(lis_PCA)-1]
plt.imshow(tr[2,:].reshape(32,32), cmap=plt.get_cmap("gray"))
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f7296e8c1d0>
```



4.1 Train feed-forward NN with reduced images

```
[ ]: def FW_iter(lis_PCA, ep, bs):  
    lis_FW = []  
    epochs = ep  
    batch_size = bs  
    for itr in range(len(lis_PCA)):  
        x_train_PCA, x_valid_PCA, x_test_PCA = lis_PCA[itr]  
  
        print("FW- components: ",(itr+1)*20-10)  
        #feed forward neural network  
        model = tf.keras.Sequential([  
            tf.keras.layers.Input(shape = (1024)),  
            tf.keras.layers.Dense(32, activation = "relu"),  
            tf.keras.layers.Dense(10, activation='softmax')  
        ])  
        model.compile(optimizer = "adam", loss='categorical_crossentropy',  
            ↪metrics=['accuracy'])  
  
        history = model.fit(x_train_PCA, y_train,  
                            batch_size = bs,
```

```

        epochs = epochs,
        validation_data=(x_valid_PCA, y_valid),
        verbose = 2

    )

    y_pred = model.predict(x_test_PCA).round()

    zo_loss = zero_one(y_pred,y_test)
    print("Zero-one loss: ",zo_loss)
    tupla = (history, model, zo_loss)
    lis_FW.append(tupla)
    return lis_FW

```

```

[ ]: epochs = 10
    batch_size = 32
    res = FW_iter(lis_PCA, epochs, batch_size)

```

```

FW- components: 10
Epoch 1/10
WARNING:tensorflow:Method (on_train_batch_end) is slow compared to the batch
update (0.175393). Check your callbacks.
1019/1019 - 5s - loss: 0.8917 - accuracy: 0.6877 - val_loss: 1.3535 -
val_accuracy: 0.5962
Epoch 2/10
1019/1019 - 5s - loss: 0.6398 - accuracy: 0.7802 - val_loss: 1.3279 -
val_accuracy: 0.6405
Epoch 3/10
1019/1019 - 5s - loss: 0.5631 - accuracy: 0.8033 - val_loss: 1.2432 -
val_accuracy: 0.6723
Epoch 4/10
1019/1019 - 5s - loss: 0.5202 - accuracy: 0.8203 - val_loss: 1.2590 -
val_accuracy: 0.6750
Epoch 5/10
1019/1019 - 5s - loss: 0.4846 - accuracy: 0.8326 - val_loss: 1.3413 -
val_accuracy: 0.6714
Epoch 6/10
1019/1019 - 5s - loss: 0.4573 - accuracy: 0.8445 - val_loss: 1.2513 -
val_accuracy: 0.6985
Epoch 7/10
1019/1019 - 5s - loss: 0.4322 - accuracy: 0.8521 - val_loss: 1.2741 -
val_accuracy: 0.6983
Epoch 8/10
1019/1019 - 5s - loss: 0.4193 - accuracy: 0.8566 - val_loss: 1.3041 -
val_accuracy: 0.7063
Epoch 9/10
1019/1019 - 5s - loss: 0.4008 - accuracy: 0.8640 - val_loss: 1.3032 -

```

```

val_accuracy: 0.7003
Epoch 10/10
1019/1019 - 5s - loss: 0.3982 - accuracy: 0.8641 - val_loss: 1.3579 -
val_accuracy: 0.6978
Zero-one loss: 0.28689275893675525
FW- components: 30
Epoch 1/10
WARNING:tensorflow:Method (on_train_batch_end) is slow compared to the batch
update (0.261349). Check your callbacks.
1019/1019 - 6s - loss: 0.6436 - accuracy: 0.7845 - val_loss: 1.0769 -
val_accuracy: 0.6961
Epoch 2/10
1019/1019 - 5s - loss: 0.3235 - accuracy: 0.8948 - val_loss: 1.0283 -
val_accuracy: 0.7299
Epoch 3/10
1019/1019 - 5s - loss: 0.2424 - accuracy: 0.9211 - val_loss: 1.0095 -
val_accuracy: 0.7387
Epoch 4/10
1019/1019 - 5s - loss: 0.1909 - accuracy: 0.9384 - val_loss: 0.9939 -
val_accuracy: 0.7633
Epoch 5/10
1019/1019 - 5s - loss: 0.1558 - accuracy: 0.9502 - val_loss: 1.1439 -
val_accuracy: 0.7466
Epoch 6/10
1019/1019 - 5s - loss: 0.1404 - accuracy: 0.9538 - val_loss: 1.0946 -
val_accuracy: 0.7713
Epoch 7/10
1019/1019 - 5s - loss: 0.1227 - accuracy: 0.9603 - val_loss: 1.1294 -
val_accuracy: 0.7809
Epoch 8/10
1019/1019 - 5s - loss: 0.1098 - accuracy: 0.9649 - val_loss: 1.1915 -
val_accuracy: 0.7785
Epoch 9/10
1019/1019 - 5s - loss: 0.1003 - accuracy: 0.9675 - val_loss: 1.1726 -
val_accuracy: 0.7894
Epoch 10/10
1019/1019 - 5s - loss: 0.0947 - accuracy: 0.9702 - val_loss: 1.2234 -
val_accuracy: 0.7822
Zero-one loss: 0.19798350137488543
FW- components: 50
Epoch 1/10
1019/1019 - 5s - loss: 0.5459 - accuracy: 0.8254 - val_loss: 0.8101 -
val_accuracy: 0.7436
Epoch 2/10
1019/1019 - 5s - loss: 0.2140 - accuracy: 0.9335 - val_loss: 0.8365 -
val_accuracy: 0.7656
Epoch 3/10
1019/1019 - 5s - loss: 0.1433 - accuracy: 0.9549 - val_loss: 0.8540 -

```

```

val_accuracy: 0.7918
Epoch 4/10
1019/1019 - 5s - loss: 0.1039 - accuracy: 0.9682 - val_loss: 0.8805 -
val_accuracy: 0.7915
Epoch 5/10
1019/1019 - 5s - loss: 0.0848 - accuracy: 0.9738 - val_loss: 0.9716 -
val_accuracy: 0.8034
Epoch 6/10
1019/1019 - 5s - loss: 0.0715 - accuracy: 0.9778 - val_loss: 0.9149 -
val_accuracy: 0.8155
Epoch 7/10
1019/1019 - 5s - loss: 0.0664 - accuracy: 0.9790 - val_loss: 1.1081 -
val_accuracy: 0.8093
Epoch 8/10
1019/1019 - 5s - loss: 0.0539 - accuracy: 0.9832 - val_loss: 1.1118 -
val_accuracy: 0.7929
Epoch 9/10
1019/1019 - 5s - loss: 0.0512 - accuracy: 0.9835 - val_loss: 1.0935 -
val_accuracy: 0.8323
Epoch 10/10
1019/1019 - 5s - loss: 0.0465 - accuracy: 0.9845 - val_loss: 1.0836 -
val_accuracy: 0.8343
Zero-one loss: 0.15902841429880843
FW- components: 70
Epoch 1/10
WARNING:tensorflow:Method (on_train_batch_end) is slow compared to the batch
update (0.200094). Check your callbacks.
1019/1019 - 5s - loss: 0.4945 - accuracy: 0.8395 - val_loss: 0.8700 -
val_accuracy: 0.7372
Epoch 2/10
1019/1019 - 5s - loss: 0.1802 - accuracy: 0.9462 - val_loss: 0.7953 -
val_accuracy: 0.7839
Epoch 3/10
1019/1019 - 5s - loss: 0.1150 - accuracy: 0.9651 - val_loss: 0.8178 -
val_accuracy: 0.8015
Epoch 4/10
1019/1019 - 5s - loss: 0.0828 - accuracy: 0.9752 - val_loss: 0.8297 -
val_accuracy: 0.7953
Epoch 5/10
1019/1019 - 5s - loss: 0.0659 - accuracy: 0.9796 - val_loss: 0.7400 -
val_accuracy: 0.8329
Epoch 6/10
1019/1019 - 5s - loss: 0.0544 - accuracy: 0.9829 - val_loss: 0.8350 -
val_accuracy: 0.8300
Epoch 7/10
1019/1019 - 5s - loss: 0.0438 - accuracy: 0.9867 - val_loss: 0.8506 -
val_accuracy: 0.8344
Epoch 8/10

```


1019/1019 - 5s - loss: 0.0406 - accuracy: 0.9879 - val_loss: 0.8914 -
 val_accuracy: 0.8357
 Epoch 9/10
 1019/1019 - 6s - loss: 0.0387 - accuracy: 0.9884 - val_loss: 0.8949 -
 val_accuracy: 0.8410
 Epoch 10/10
 1019/1019 - 6s - loss: 0.0312 - accuracy: 0.9904 - val_loss: 0.9068 -
 val_accuracy: 0.8454
 Zero-one loss: 0.1457378551787351
 FW- components: 90
 Epoch 1/10
 1019/1019 - 5s - loss: 0.5080 - accuracy: 0.8416 - val_loss: 0.7913 -
 val_accuracy: 0.7591
 Epoch 2/10
 1019/1019 - 5s - loss: 0.1627 - accuracy: 0.9520 - val_loss: 0.7948 -
 val_accuracy: 0.7869
 Epoch 3/10
 1019/1019 - 5s - loss: 0.0948 - accuracy: 0.9742 - val_loss: 0.7662 -
 val_accuracy: 0.8137
 Epoch 4/10
 1019/1019 - 5s - loss: 0.0640 - accuracy: 0.9823 - val_loss: 0.7883 -
 val_accuracy: 0.8290
 Epoch 5/10
 1019/1019 - 5s - loss: 0.0502 - accuracy: 0.9860 - val_loss: 0.7893 -
 val_accuracy: 0.8356
 Epoch 6/10
 1019/1019 - 5s - loss: 0.0395 - accuracy: 0.9891 - val_loss: 0.8706 -
 val_accuracy: 0.8399
 Epoch 7/10
 1019/1019 - 5s - loss: 0.0400 - accuracy: 0.9881 - val_loss: 1.0676 -
 val_accuracy: 0.8155
 Epoch 8/10
 1019/1019 - 5s - loss: 0.0282 - accuracy: 0.9915 - val_loss: 1.0146 -
 val_accuracy: 0.8299
 Epoch 9/10
 1019/1019 - 5s - loss: 0.0238 - accuracy: 0.9928 - val_loss: 0.9587 -
 val_accuracy: 0.8548
 Epoch 10/10
 1019/1019 - 5s - loss: 0.0224 - accuracy: 0.9929 - val_loss: 1.0540 -
 val_accuracy: 0.8471
 Zero-one loss: 0.15627864344637946
 FW- components: 110
 Epoch 1/10
 1019/1019 - 5s - loss: 0.4724 - accuracy: 0.8546 - val_loss: 0.8186 -
 val_accuracy: 0.7603
 Epoch 2/10
 1019/1019 - 5s - loss: 0.1382 - accuracy: 0.9634 - val_loss: 0.7725 -
 val_accuracy: 0.8089

Epoch 3/10
1019/1019 - 5s - loss: 0.0794 - accuracy: 0.9783 - val_loss: 0.8203 - val_accuracy: 0.8191

Epoch 4/10
1019/1019 - 5s - loss: 0.0528 - accuracy: 0.9850 - val_loss: 0.9337 - val_accuracy: 0.8204

Epoch 5/10
1019/1019 - 5s - loss: 0.0417 - accuracy: 0.9879 - val_loss: 0.9542 - val_accuracy: 0.8313

Epoch 6/10
1019/1019 - 5s - loss: 0.0369 - accuracy: 0.9895 - val_loss: 0.9346 - val_accuracy: 0.8412

Epoch 7/10
1019/1019 - 5s - loss: 0.0267 - accuracy: 0.9927 - val_loss: 1.0282 - val_accuracy: 0.8474

Epoch 8/10
1019/1019 - 5s - loss: 0.0211 - accuracy: 0.9939 - val_loss: 1.1628 - val_accuracy: 0.8286

Epoch 9/10
1019/1019 - 5s - loss: 0.0264 - accuracy: 0.9914 - val_loss: 1.1593 - val_accuracy: 0.8408

Epoch 10/10
1019/1019 - 5s - loss: 0.0205 - accuracy: 0.9938 - val_loss: 1.2966 - val_accuracy: 0.8436

Zero-one loss: 0.14848762603116408

FW- components: 130

Epoch 1/10
WARNING:tensorflow:Method (on_train_batch_end) is slow compared to the batch update (0.148986). Check your callbacks.

1019/1019 - 5s - loss: 0.5150 - accuracy: 0.8395 - val_loss: 0.7711 - val_accuracy: 0.7811

Epoch 2/10
1019/1019 - 5s - loss: 0.1556 - accuracy: 0.9550 - val_loss: 0.7043 - val_accuracy: 0.8064

Epoch 3/10
1019/1019 - 5s - loss: 0.0860 - accuracy: 0.9749 - val_loss: 0.7319 - val_accuracy: 0.8285

Epoch 4/10
1019/1019 - 5s - loss: 0.0547 - accuracy: 0.9850 - val_loss: 0.8580 - val_accuracy: 0.8292

Epoch 5/10
1019/1019 - 5s - loss: 0.0448 - accuracy: 0.9874 - val_loss: 0.9114 - val_accuracy: 0.8353

Epoch 6/10
1019/1019 - 5s - loss: 0.0326 - accuracy: 0.9913 - val_loss: 0.9489 - val_accuracy: 0.8368

Epoch 7/10
1019/1019 - 5s - loss: 0.0302 - accuracy: 0.9910 - val_loss: 0.8936 -

```

val_accuracy: 0.8608
Epoch 8/10
1019/1019 - 5s - loss: 0.0265 - accuracy: 0.9921 - val_loss: 0.9696 -
val_accuracy: 0.8584
Epoch 9/10
1019/1019 - 5s - loss: 0.0264 - accuracy: 0.9920 - val_loss: 1.0818 -
val_accuracy: 0.8422
Epoch 10/10
1019/1019 - 5s - loss: 0.0166 - accuracy: 0.9955 - val_loss: 1.1559 -
val_accuracy: 0.8453
Zero-one loss: 0.152153987167736
FW- components: 150
Epoch 1/10
1019/1019 - 5s - loss: 0.4667 - accuracy: 0.8549 - val_loss: 0.7300 -
val_accuracy: 0.7691
Epoch 2/10
1019/1019 - 5s - loss: 0.1365 - accuracy: 0.9635 - val_loss: 0.6678 -
val_accuracy: 0.8116
Epoch 3/10
1019/1019 - 5s - loss: 0.0762 - accuracy: 0.9800 - val_loss: 0.6308 -
val_accuracy: 0.8370
Epoch 4/10
1019/1019 - 5s - loss: 0.0458 - accuracy: 0.9884 - val_loss: 0.6961 -
val_accuracy: 0.8396
Epoch 5/10
1019/1019 - 5s - loss: 0.0426 - accuracy: 0.9883 - val_loss: 0.8004 -
val_accuracy: 0.8357
Epoch 6/10
1019/1019 - 5s - loss: 0.0267 - accuracy: 0.9929 - val_loss: 0.7318 -
val_accuracy: 0.8495
Epoch 7/10
1019/1019 - 5s - loss: 0.0267 - accuracy: 0.9921 - val_loss: 0.7954 -
val_accuracy: 0.8433
Epoch 8/10
1019/1019 - 5s - loss: 0.0232 - accuracy: 0.9931 - val_loss: 1.0402 -
val_accuracy: 0.8232
Epoch 9/10
1019/1019 - 5s - loss: 0.0190 - accuracy: 0.9945 - val_loss: 0.9076 -
val_accuracy: 0.8549
Epoch 10/10
1019/1019 - 5s - loss: 0.0184 - accuracy: 0.9949 - val_loss: 0.9122 -
val_accuracy: 0.8505
Zero-one loss: 0.14115490375802017
FW- components: 170
Epoch 1/10
1019/1019 - 5s - loss: 0.5013 - accuracy: 0.8487 - val_loss: 0.7488 -
val_accuracy: 0.7780
Epoch 2/10

```

1019/1019 - 5s - loss: 0.1291 - accuracy: 0.9651 - val_loss: 0.7400 -
 val_accuracy: 0.8058
 Epoch 3/10
 1019/1019 - 5s - loss: 0.0664 - accuracy: 0.9830 - val_loss: 0.6988 -
 val_accuracy: 0.8416
 Epoch 4/10
 1019/1019 - 5s - loss: 0.0447 - accuracy: 0.9883 - val_loss: 0.7572 -
 val_accuracy: 0.8386
 Epoch 5/10
 1019/1019 - 5s - loss: 0.0339 - accuracy: 0.9907 - val_loss: 0.8077 -
 val_accuracy: 0.8454
 Epoch 6/10
 1019/1019 - 5s - loss: 0.0254 - accuracy: 0.9930 - val_loss: 0.7187 -
 val_accuracy: 0.8663
 Epoch 7/10
 1019/1019 - 5s - loss: 0.0213 - accuracy: 0.9939 - val_loss: 0.7830 -
 val_accuracy: 0.8504
 Epoch 8/10
 1019/1019 - 5s - loss: 0.0189 - accuracy: 0.9946 - val_loss: 0.7266 -
 val_accuracy: 0.8778
 Epoch 9/10
 1019/1019 - 5s - loss: 0.0124 - accuracy: 0.9968 - val_loss: 0.8611 -
 val_accuracy: 0.8610
 Epoch 10/10
 1019/1019 - 5s - loss: 0.0178 - accuracy: 0.9946 - val_loss: 0.8310 -
 val_accuracy: 0.8737
 Zero-one loss: 0.12465627864344637
 FW- components: 190
 Epoch 1/10
 1019/1019 - 5s - loss: 0.4545 - accuracy: 0.8607 - val_loss: 0.7529 -
 val_accuracy: 0.7865
 Epoch 2/10
 1019/1019 - 5s - loss: 0.1264 - accuracy: 0.9652 - val_loss: 0.7090 -
 val_accuracy: 0.8212
 Epoch 3/10
 1019/1019 - 5s - loss: 0.0705 - accuracy: 0.9815 - val_loss: 0.6724 -
 val_accuracy: 0.8435
 Epoch 4/10
 1019/1019 - 5s - loss: 0.0461 - accuracy: 0.9876 - val_loss: 0.7778 -
 val_accuracy: 0.8344
 Epoch 5/10
 1019/1019 - 5s - loss: 0.0383 - accuracy: 0.9895 - val_loss: 0.7378 -
 val_accuracy: 0.8549
 Epoch 6/10
 1019/1019 - 5s - loss: 0.0264 - accuracy: 0.9932 - val_loss: 0.9976 -
 val_accuracy: 0.8267
 Epoch 7/10
 1019/1019 - 5s - loss: 0.0244 - accuracy: 0.9929 - val_loss: 0.9187 -

```

val_accuracy: 0.8483
Epoch 8/10
1019/1019 - 5s - loss: 0.0181 - accuracy: 0.9950 - val_loss: 0.9842 -
val_accuracy: 0.8478
Epoch 9/10
1019/1019 - 6s - loss: 0.0234 - accuracy: 0.9936 - val_loss: 0.9568 -
val_accuracy: 0.8530
Epoch 10/10
1019/1019 - 6s - loss: 0.0140 - accuracy: 0.9960 - val_loss: 0.9279 -
val_accuracy: 0.8657
Zero-one loss: 0.12557286892758937
FW- components: 210
Epoch 1/10
1019/1019 - 5s - loss: 0.4542 - accuracy: 0.8597 - val_loss: 0.7646 -
val_accuracy: 0.7775
Epoch 2/10
1019/1019 - 5s - loss: 0.1189 - accuracy: 0.9685 - val_loss: 0.7023 -
val_accuracy: 0.8039
Epoch 3/10
1019/1019 - 5s - loss: 0.0641 - accuracy: 0.9839 - val_loss: 0.7556 -
val_accuracy: 0.8193
Epoch 4/10
1019/1019 - 5s - loss: 0.0412 - accuracy: 0.9896 - val_loss: 0.8243 -
val_accuracy: 0.8222
Epoch 5/10
1019/1019 - 5s - loss: 0.0262 - accuracy: 0.9936 - val_loss: 0.7042 -
val_accuracy: 0.8571
Epoch 6/10
1019/1019 - 5s - loss: 0.0262 - accuracy: 0.9922 - val_loss: 0.7838 -
val_accuracy: 0.8453
Epoch 7/10
1019/1019 - 5s - loss: 0.0157 - accuracy: 0.9956 - val_loss: 0.8158 -
val_accuracy: 0.8592
Epoch 8/10
1019/1019 - 5s - loss: 0.0245 - accuracy: 0.9926 - val_loss: 0.9040 -
val_accuracy: 0.8639
Epoch 9/10
1019/1019 - 5s - loss: 0.0147 - accuracy: 0.9961 - val_loss: 0.9469 -
val_accuracy: 0.8372
Epoch 10/10
1019/1019 - 5s - loss: 0.0174 - accuracy: 0.9947 - val_loss: 1.0241 -
val_accuracy: 0.8486
Zero-one loss: 0.1457378551787351

```

```

[ ]: df = pd.DataFrame(
      columns = ['epochs', 'valid', 'components', 'accuracy', 'value']
    )

```

```

for itr in range(len(res)):

    time = [i for i in range(1,epochs+1)]
    valids = [0 for i in range(1,epochs+1)]
    components = [(itr+1)*20-10 for i in range(1,epochs+1)]

    accur = [1 for i in range(1,epochs+1)]
    acc = res[itr][0].history['accuracy']

    df1= pd.DataFrame(data= np.vstack((time,valids,components,accur,acc)).T,
    →columns = ['epochs', 'valid', 'components', 'accuracy', 'value'])

    loss= res[itr][0].history['loss']
    accur = [0 for i in range(1,epochs+1)]
    df2= pd.DataFrame(data= np.vstack((time,valids,components,accur,loss)).T,
    →columns = ['epochs', 'valid', 'components', 'accuracy', 'value'])

    valids = [1 for i in range(1,epochs+1)]
    accur = [1 for i in range(1,epochs+1)]
    val_acc = res[itr][0].history['val_accuracy']

    df3= pd.DataFrame(data= np.vstack((time,valids,components,accur,val_acc)).
    →T, columns = ['epochs', 'valid', 'components', 'accuracy', 'value'])

    accur = [0 for i in range(1,epochs+1)]

    val_loss = res[itr][0].history['val_loss']
    df4= pd.DataFrame(data= np.vstack((time,valids,components,accur,val_loss)).
    →T, columns = ['epochs', 'valid', 'components', 'accuracy', 'value'])

    df = df.append(df1.append(df2).append(df3).append(df4))

```

```

[ ]: df['components'] = df['components'].astype('category')
df = df.assign(accuracy = ['accuracy' if accuracy == 1. else 'loss' for
    →accuracy in df['accuracy']])
df = df.assign(valid = ['validation' if valid == 1. else 'training' for valid
    →in df['valid']])
df['accuracy'].unique()

```

```

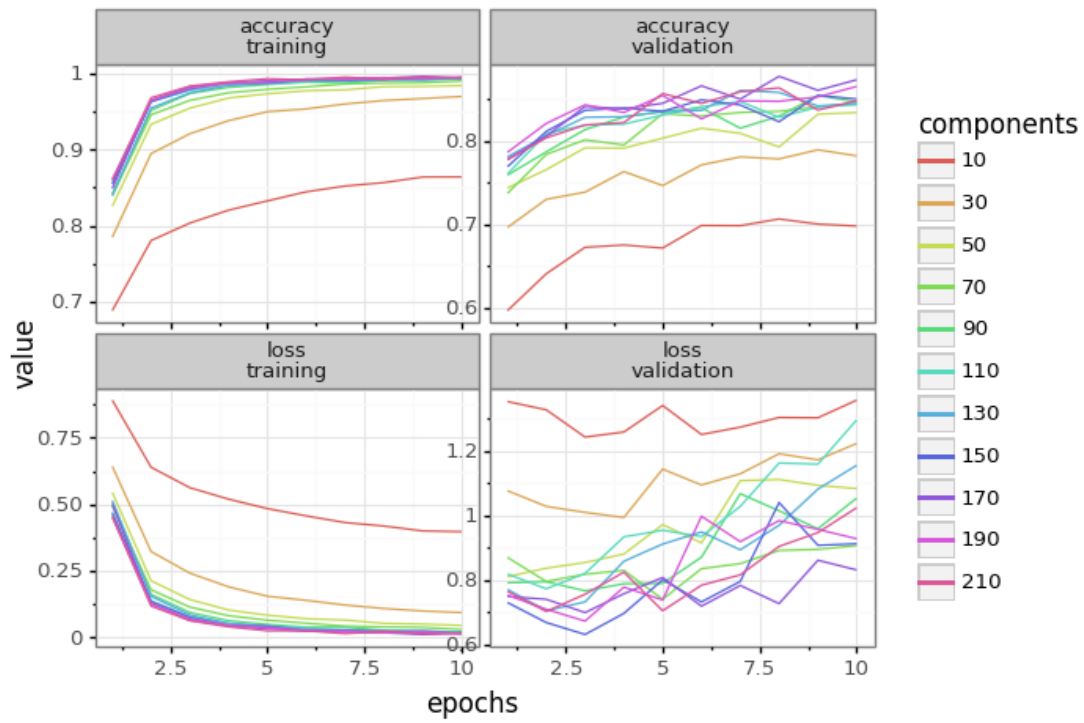
[ ]: array(['accuracy', 'loss'], dtype=object)

```

```

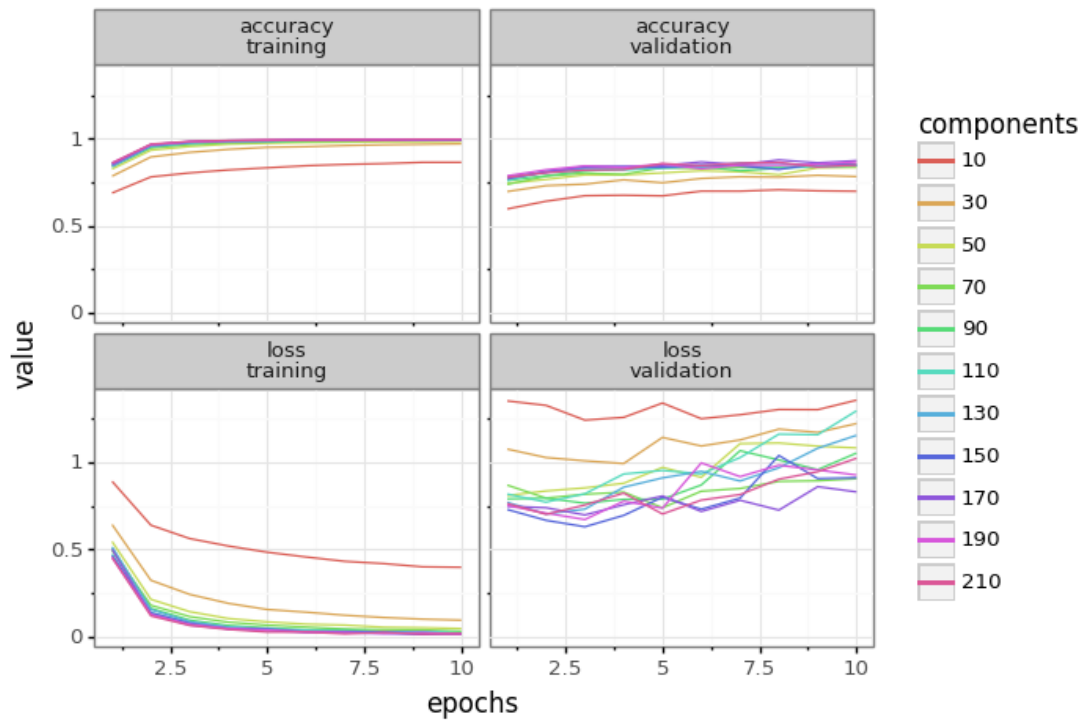
[ ]: ggplot(df, aes(x='epochs', y='value',color='components')) + \
    geom_line() + \
    facet_wrap(['accuracy', 'valid'],scales='free') + theme_bw(base_size=12)

```



```
[ ]: <ggplot: (8758028328796)>
```

```
[ ]: ggplot(df, aes(x='epochs', y='value', color='components')) + \
      geom_line() + \
      facet_wrap(['accuracy', 'valid']) + \
      theme_bw(base_size=12)
```



```
[ ]: <ggplot: (-9223363278826112789)>
```

4.2 Feed-Forward Zero-One Loss

```
[ ]: df_loss = pd.DataFrame(
    columns = ['components', 'zero_one']
)

losses = []

for i in range(len(res)):
    losses.append(res[i][2])
components = [i for i in range(10,211,20)]
```


```
[ ]: df_loss = pd.DataFrame( data = [components,losses], index = ['components', 'zero_one']).T
```

```
[ ]: df_loss
```

```
[ ]:
   components  zero_one
0         10.0   0.286893
1         30.0   0.197984
2         50.0   0.159028
```


3	70.0	0.145738
4	90.0	0.156279
5	110.0	0.148488
6	130.0	0.152154
7	150.0	0.141155
8	170.0	0.124656
9	190.0	0.125573
10	210.0	0.145738

```
[ ]: ggplot(df_loss, aes(x='components', y='zero_one')) + \
  geom_line() + \
  geom_point() + \
  theme_bw(base_size=12) + ggtitle("Zero-one loss FW") + ylab("loss")
```



```
[ ]: <ggplot: (-9223363278721333202)>
```

5 1.4 Convolutional Neural Newtworks

5.1 One VGG block CNN

```
[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu',
    ↪input_shape = (32, 32, 3)),
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation = 'relu'),
    tf.keras.layers.Dense(10, activation = 'softmax')
])

model.compile(optimizer = "adam", loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

model.summary()
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
flatten (Flatten)	(None, 8192)	0
dense_22 (Dense)	(None, 128)	1048704
dense_23 (Dense)	(None, 10)	1290
Total params: 1,060,138		
Trainable params: 1,060,138		
Non-trainable params: 0		

Model: "sequential_11"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0

```

-----
flatten (Flatten)                (None, 8192)                0
-----
dense_22 (Dense)                 (None, 128)                1048704
-----
dense_23 (Dense)                 (None, 10)                 1290
=====
Total params: 1,060,138
Trainable params: 1,060,138
Non-trainable params: 0
-----

```

```

[ ]: history = model.fit(x_train,y_train,
    batch_size = 32,
    epochs=10,
    validation_data=(x_valid, y_valid),
    verbose=2
    )

```

```

Epoch 1/10
1019/1019 - 12s - loss: 0.2258 - accuracy: 0.9251 - val_loss: 0.1457 -
val_accuracy: 0.9591
Epoch 2/10
1019/1019 - 11s - loss: 0.0122 - accuracy: 0.9964 - val_loss: 0.1096 -
val_accuracy: 0.9716
Epoch 3/10
1019/1019 - 11s - loss: 0.0083 - accuracy: 0.9981 - val_loss: 0.0620 -
val_accuracy: 0.9841
Epoch 4/10
1019/1019 - 11s - loss: 0.0058 - accuracy: 0.9989 - val_loss: 0.1105 -
val_accuracy: 0.9822
Epoch 5/10
1019/1019 - 11s - loss: 0.0071 - accuracy: 0.9986 - val_loss: 0.1179 -
val_accuracy: 0.9836
Epoch 6/10
1019/1019 - 11s - loss: 0.0080 - accuracy: 0.9989 - val_loss: 0.1146 -
val_accuracy: 0.9782
Epoch 7/10
1019/1019 - 11s - loss: 0.0039 - accuracy: 0.9994 - val_loss: 0.1395 -
val_accuracy: 0.9790
Epoch 8/10
1019/1019 - 11s - loss: 0.0020 - accuracy: 0.9995 - val_loss: 0.1240 -
val_accuracy: 0.9827
Epoch 9/10
1019/1019 - 11s - loss: 0.0045 - accuracy: 0.9995 - val_loss: 0.1455 -
val_accuracy: 0.9733
Epoch 10/10
1019/1019 - 11s - loss: 0.0029 - accuracy: 0.9994 - val_loss: 0.1433 -

```

val_accuracy: 0.9771


```
[ ]: y_pred = model.predict(x_test)

cnn_loss = []
zol = zero_one(y_pred, y_test)

print("Zero-one Loss: ", zol)
cnn_loss.append(zol)
```

Zero-one Loss: 0.02153987167736022

```
[ ]: # plot a random sample of test images, their predicted labels, and ground truth
fig = plt.figure(figsize=(16, 9))
for i, idx in enumerate(np.random.choice(x_test.shape[0], size=16,
↪replace=False)):
    ax = fig.add_subplot(4, 4, i + 1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(x_test[idx]))
    pred_idx = np.argmax(y_pred[idx])
    true_idx = np.argmax(y_test[idx])
    ax.set_title("{} ({}).format(TYPES[pred_idx], TYPES[true_idx]),
                color=("green" if pred_idx == true_idx else "red"))
```




```
[ ]: #Loss and accuracy visualisation

plt.figure(1)

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
[ ]:
```

5.2 Two VGG blocks

```
[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu',
    ↪input_shape = (32, 32, 3)),
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.Conv2D(64, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation = 'relu'),
    tf.keras.layers.Dense(10, activation = 'softmax')
])

model.compile(optimizer = "adam", loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

model.summary()
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 32, 32, 32)	896
conv2d_3 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_5 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten_1 (Flatten)	(None, 4096)	0
dense_24 (Dense)	(None, 128)	524416
dense_25 (Dense)	(None, 10)	1290

=====
Total params: 591,274
Trainable params: 591,274
Non-trainable params: 0

```
-----  
[ ]: history2 = model.fit(x_train,y_train,  
    batch_size = 32,  
    epochs=10,  
    validation_data=(x_valid, y_valid),  
    verbose=2  
  
    )
```

Epoch 1/10

WARNING:tensorflow:Method (on_train_batch_end) is slow compared to the batch update (0.210605). Check your callbacks.

1019/1019 - 14s - loss: 0.2516 - accuracy: 0.9158 - val_loss: 0.1399 - val_accuracy: 0.9559

Epoch 2/10

1019/1019 - 13s - loss: 0.0240 - accuracy: 0.9931 - val_loss: 0.0467 - val_accuracy: 0.9812

Epoch 3/10

1019/1019 - 13s - loss: 0.0077 - accuracy: 0.9973 - val_loss: 0.0338 - val_accuracy: 0.9900

Epoch 4/10

1019/1019 - 13s - loss: 3.4333e-05 - accuracy: 1.0000 - val_loss: 0.0311 - val_accuracy: 0.9900

Epoch 5/10

1019/1019 - 13s - loss: 1.1592e-05 - accuracy: 1.0000 - val_loss: 0.0367 - val_accuracy: 0.9900

Epoch 6/10

1019/1019 - 13s - loss: 5.8665e-06 - accuracy: 1.0000 - val_loss: 0.0342 - val_accuracy: 0.9903

Epoch 7/10

1019/1019 - 13s - loss: 2.9794e-06 - accuracy: 1.0000 - val_loss: 0.0304 - val_accuracy: 0.9903

Epoch 8/10

1019/1019 - 13s - loss: 1.6845e-06 - accuracy: 1.0000 - val_loss: 0.0328 - val_accuracy: 0.9906

Epoch 9/10

1019/1019 - 13s - loss: 9.3372e-07 - accuracy: 1.0000 - val_loss: 0.0308 - val_accuracy: 0.9905

Epoch 10/10

1019/1019 - 14s - loss: 5.2326e-07 - accuracy: 1.0000 - val_loss: 0.0337 - val_accuracy: 0.9907

```
[ ]: # evaluate zero-one loss  
y_pred = model.predict(x_test)  
zol = zero_one(y_pred, y_test)
```

```
print("Zero-one Loss: ", zol)

cnn_loss.append(zol)
```

Zero-one Loss: 0.00916590284142988

5.3 Three VGG blocks

```
[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu',
    ↪input_shape = (32, 32, 3)),
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.Conv2D(64, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Conv2D(128, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.Conv2D(128, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation = 'relu'),
    tf.keras.layers.Dense(10, activation = 'softmax')
])

model.compile(optimizer = "adam", loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

model.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 32)	896
conv2d_7 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_3 (MaxPooling2)	(None, 16, 16, 32)	0
conv2d_8 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_9 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_4 (MaxPooling2)	(None, 8, 8, 64)	0
conv2d_10 (Conv2D)	(None, 8, 8, 128)	73856

conv2d_11 (Conv2D)	(None, 8, 8, 128)	147584

max_pooling2d_5 (MaxPooling2)	(None, 4, 4, 128)	0

flatten_2 (Flatten)	(None, 2048)	0

dense_26 (Dense)	(None, 128)	262272

dense_27 (Dense)	(None, 10)	1290
=====		
Total params: 550,570		
Trainable params: 550,570		
Non-trainable params: 0		

```
[ ]: history3 = model.fit(x_train,y_train,
    batch_size = 32,
    epochs=10,
    validation_data=(x_valid, y_valid),
    verbose=2

    )
```

```
Epoch 1/10
1019/1019 - 19s - loss: 0.3207 - accuracy: 0.8886 - val_loss: 0.0736 -
val_accuracy: 0.9760
Epoch 2/10
1019/1019 - 18s - loss: 0.0340 - accuracy: 0.9895 - val_loss: 0.2608 -
val_accuracy: 0.9379
Epoch 3/10
1019/1019 - 18s - loss: 0.0234 - accuracy: 0.9932 - val_loss: 0.0518 -
val_accuracy: 0.9861
Epoch 4/10
1019/1019 - 18s - loss: 8.6958e-05 - accuracy: 1.0000 - val_loss: 0.0193 -
val_accuracy: 0.9917
Epoch 5/10
1019/1019 - 18s - loss: 9.9077e-06 - accuracy: 1.0000 - val_loss: 0.0171 -
val_accuracy: 0.9928
Epoch 6/10
1019/1019 - 18s - loss: 4.4342e-06 - accuracy: 1.0000 - val_loss: 0.0164 -
val_accuracy: 0.9925
Epoch 7/10
1019/1019 - 18s - loss: 2.2970e-06 - accuracy: 1.0000 - val_loss: 0.0168 -
val_accuracy: 0.9925
Epoch 8/10
1019/1019 - 18s - loss: 1.2628e-06 - accuracy: 1.0000 - val_loss: 0.0164 -
val_accuracy: 0.9932
Epoch 9/10
```

```
1019/1019 - 18s - loss: 6.9922e-07 - accuracy: 1.0000 - val_loss: 0.0156 -
val_accuracy: 0.9940
Epoch 10/10
1019/1019 - 18s - loss: 3.8787e-07 - accuracy: 1.0000 - val_loss: 0.0166 -
val_accuracy: 0.9938
```

```
[ ]: # evaluate zero-one loss
y_pred = model.predict(x_test)

zol = zero_one(y_pred, y_test)

print("Zero-one Loss: ", zol)

cnn_loss.append(zol)
```

Zero-one Loss: 0.00916590284142988

5.4 Three VGG blocks with Dropout

```
[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu',
    ↪input_shape = (32, 32, 3)),
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Conv2D(64, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.Conv2D(64, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Conv2D(128, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.Conv2D(128, (3, 3), padding = "same", activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation = 'relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation = 'softmax')
])

model.compile(optimizer = "adam", loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

model.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
conv2d_12 (Conv2D)          (None, 32, 32, 32)      896
-----
conv2d_13 (Conv2D)          (None, 32, 32, 32)      9248
-----
max_pooling2d_6 (MaxPooling2 (None, 16, 16, 32)      0
-----
dropout (Dropout)          (None, 16, 16, 32)      0
-----
conv2d_14 (Conv2D)          (None, 16, 16, 64)      18496
-----
conv2d_15 (Conv2D)          (None, 16, 16, 64)      36928
-----
max_pooling2d_7 (MaxPooling2 (None, 8, 8, 64)      0
-----
dropout_1 (Dropout)         (None, 8, 8, 64)      0
-----
conv2d_16 (Conv2D)          (None, 8, 8, 128)      73856
-----
conv2d_17 (Conv2D)          (None, 8, 8, 128)      147584
-----
max_pooling2d_8 (MaxPooling2 (None, 4, 4, 128)      0
-----
dropout_2 (Dropout)         (None, 4, 4, 128)      0
-----
flatten_3 (Flatten)         (None, 2048)            0
-----
dense_28 (Dense)            (None, 128)             262272
-----
dropout_3 (Dropout)         (None, 128)             0
-----
dense_29 (Dense)            (None, 10)              1290
=====
Total params: 550,570
Trainable params: 550,570
Non-trainable params: 0

```

```

[ ]: history4 = model.fit(x_train,y_train,
    batch_size = 32,
    epochs=10,
    validation_data=(x_valid, y_valid),
    verbose=2
    )

```

Epoch 1/10

WARNING:tensorflow:Method (on_train_batch_end) is slow compared to the batch update (0.151422). Check your callbacks.

```

1019/1019 - 20s - loss: 0.4977 - accuracy: 0.8243 - val_loss: 0.1280 -
val_accuracy: 0.9546
Epoch 2/10
1019/1019 - 19s - loss: 0.0688 - accuracy: 0.9775 - val_loss: 0.1614 -
val_accuracy: 0.9452
Epoch 3/10
1019/1019 - 19s - loss: 0.0458 - accuracy: 0.9857 - val_loss: 0.0644 -
val_accuracy: 0.9795
Epoch 4/10
1019/1019 - 19s - loss: 0.0364 - accuracy: 0.9886 - val_loss: 0.0873 -
val_accuracy: 0.9727
Epoch 5/10
1019/1019 - 19s - loss: 0.0243 - accuracy: 0.9931 - val_loss: 0.1079 -
val_accuracy: 0.9783
Epoch 6/10
1019/1019 - 19s - loss: 0.0328 - accuracy: 0.9902 - val_loss: 0.0799 -
val_accuracy: 0.9782
Epoch 7/10
1019/1019 - 20s - loss: 0.0262 - accuracy: 0.9922 - val_loss: 0.2361 -
val_accuracy: 0.9506
Epoch 8/10
1019/1019 - 19s - loss: 0.0189 - accuracy: 0.9949 - val_loss: 0.1377 -
val_accuracy: 0.9713
Epoch 9/10
1019/1019 - 19s - loss: 0.0246 - accuracy: 0.9928 - val_loss: 0.0759 -
val_accuracy: 0.9793
Epoch 10/10
1019/1019 - 19s - loss: 0.0190 - accuracy: 0.9952 - val_loss: 0.3236 -
val_accuracy: 0.9211

```

```

[ ]: # evaluate zero-one loss
y_pred = model.predict(x_test)
zol = zero_one(y_pred, y_test)

print("Zero-one Loss: ", zol)

cnn_loss.append(zol)

```

Zero-one Loss: 0.07561869844179651

5.5 Three VGG blocks with Dropout and Batch Normalization

```

[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), padding = "same", use_bias=False,
    ↪input_shape = (32, 32, 3)),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Activation('relu'),

```

```

tf.keras.layers.Conv2D(32, (3, 3), padding = "same", use_bias=False),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Activation('relu'),
tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
tf.keras.layers.Dropout(0.2),
tf.keras.layers.Conv2D(64, (3, 3), padding = "same", use_bias=False),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Activation('relu'),
tf.keras.layers.Conv2D(64, (3, 3), padding = "same", use_bias=False),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Activation('relu'),
tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
tf.keras.layers.Dropout(0.3),
tf.keras.layers.Conv2D(128, (3, 3), padding = "same", use_bias=False),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Activation('relu'),
tf.keras.layers.Conv2D(128, (3, 3), padding = "same", use_bias=False),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Activation('relu'),
tf.keras.layers.MaxPooling2D(pool_size = (2, 2)),
tf.keras.layers.Dropout(0.4),
tf.keras.layers.Flatten(),
tf.keras.layers.Dense(128, use_bias=False),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Activation('relu'),
tf.keras.layers.Dropout(0.5),
tf.keras.layers.Dense(10, activation = 'softmax')
])

model.compile(optimizer = "adam", loss='categorical_crossentropy',
↳metrics=['accuracy'])

model.summary()

```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 32, 32, 32)	864
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
activation (Activation)	(None, 32, 32, 32)	0
conv2d_19 (Conv2D)	(None, 32, 32, 32)	9216
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128

activation_1 (Activation)	(None, 32, 32, 32)	0
max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_4 (Dropout)	(None, 16, 16, 32)	0
conv2d_20 (Conv2D)	(None, 16, 16, 64)	18432
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 64)	256
activation_2 (Activation)	(None, 16, 16, 64)	0
conv2d_21 (Conv2D)	(None, 16, 16, 64)	36864
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 64)	256
activation_3 (Activation)	(None, 16, 16, 64)	0
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_5 (Dropout)	(None, 8, 8, 64)	0
conv2d_22 (Conv2D)	(None, 8, 8, 128)	73728
batch_normalization_4 (Batch Normalization)	(None, 8, 8, 128)	512
activation_4 (Activation)	(None, 8, 8, 128)	0
conv2d_23 (Conv2D)	(None, 8, 8, 128)	147456
batch_normalization_5 (Batch Normalization)	(None, 8, 8, 128)	512
activation_5 (Activation)	(None, 8, 8, 128)	0
max_pooling2d_11 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_6 (Dropout)	(None, 4, 4, 128)	0
flatten_4 (Flatten)	(None, 2048)	0
dense_30 (Dense)	(None, 128)	262144
batch_normalization_6 (Batch Normalization)	(None, 128)	512
activation_6 (Activation)	(None, 128)	0
dropout_7 (Dropout)	(None, 128)	0

```

-----
dense_31 (Dense)                (None, 10)                1290
=====
Total params: 552,298
Trainable params: 551,146
Non-trainable params: 1,152
-----

```

```
[ ]:
```

```
[ ]: history5 = model.fit(x_train,y_train,
    batch_size = 32,
    epochs=10,
    validation_data=(x_valid, y_valid),
    verbose=2
    )
```

```

Epoch 1/10
1019/1019 - 22s - loss: 0.3911 - accuracy: 0.8740 - val_loss: 0.1592 -
val_accuracy: 0.9462
Epoch 2/10
1019/1019 - 21s - loss: 0.0870 - accuracy: 0.9735 - val_loss: 0.0608 -
val_accuracy: 0.9725
Epoch 3/10
1019/1019 - 21s - loss: 0.0421 - accuracy: 0.9875 - val_loss: 0.0188 -
val_accuracy: 0.9944
Epoch 4/10
1019/1019 - 21s - loss: 0.0428 - accuracy: 0.9876 - val_loss: 0.0137 -
val_accuracy: 0.9958
Epoch 5/10
1019/1019 - 21s - loss: 0.0259 - accuracy: 0.9927 - val_loss: 0.0061 -
val_accuracy: 0.9978
Epoch 6/10
1019/1019 - 22s - loss: 0.0252 - accuracy: 0.9924 - val_loss: 0.0168 -
val_accuracy: 0.9929
Epoch 7/10
1019/1019 - 21s - loss: 0.0234 - accuracy: 0.9928 - val_loss: 0.0430 -
val_accuracy: 0.9826
Epoch 8/10
1019/1019 - 21s - loss: 0.0239 - accuracy: 0.9925 - val_loss: 0.0387 -
val_accuracy: 0.9829
Epoch 9/10
1019/1019 - 21s - loss: 0.0169 - accuracy: 0.9948 - val_loss: 0.0207 -
val_accuracy: 0.9916
Epoch 10/10
1019/1019 - 21s - loss: 0.0140 - accuracy: 0.9956 - val_loss: 0.0163 -
val_accuracy: 0.9936

```

```
[ ]: # evaluate zero-one loss
y_pred = model.predict(x_test)
zol = zero_one(y_pred, y_test)

print("Zero-one Loss: ", zol)

cnn_loss.append(zol)
```

Zero-one Loss: 0.005957836846929423


5.6 VGG CNN results

One VGG block

```
[ ]: plt.figure(1)

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```





Two VGG block

```
[ ]: plt.figure(1)

plt.subplot(211)
plt.plot(history2.history['accuracy'])
plt.plot(history2.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history2.history['loss'])
plt.plot(history2.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```




Three VGG block

```
[ ]: plt.figure(1)

plt.subplot(211)
plt.plot(history3.history['accuracy'])
plt.plot(history3.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history3.history['loss'])
plt.plot(history3.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```




Three VGG block with Dropout

```
[ ]: plt.figure(1)

plt.subplot(211)
plt.plot(history4.history['accuracy'])
plt.plot(history4.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history4.history['loss'])
plt.plot(history4.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

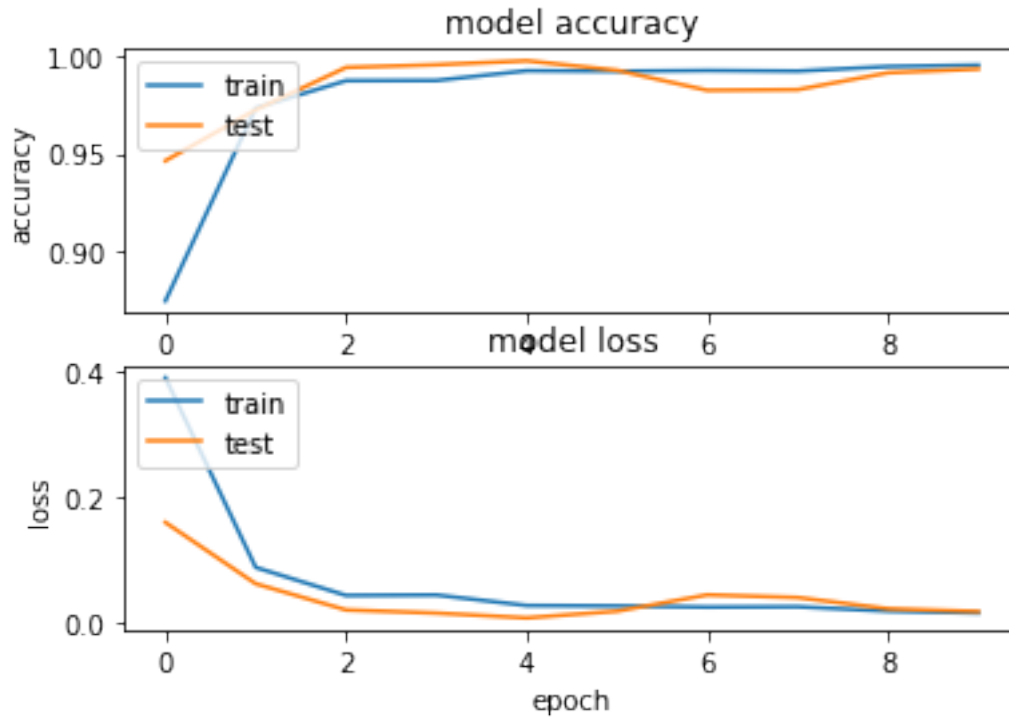


Three VGG block with Dropout and Batch Normalization

```
[ ]: plt.figure(1)

plt.subplot(211)
plt.plot(history5.history['accuracy'])
plt.plot(history5.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history5.history['loss'])
plt.plot(history5.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



5.7 VGG Zero-One Loss

```
[ ]: cnn_loss
```

```
[ ]: [0.02153987167736022,
      0.00916590284142988,
      0.00916590284142988,
      0.07561869844179651,
      0.005957836846929423]
```

```
[ ]: df_cnn_loss = pd.DataFrame(
      columns = ['type', 'zero_one']
    )

      types = ["1VGG","2VGG","3VGG","3VGG-drop", "3VGG-drop-norm"]
      df_cnn_loss = pd.DataFrame( data = [types,cnn_loss], index = ['types',
      ↪ 'zero_one']).T
```

```
[ ]: df_cnn_loss
```

```
[ ]:      types      zero_one
      0      1VGG      0.021539
      1      2VGG      0.009165
```

```


2          3VGG      0.0091659
3      3VGG-drop    0.0756187
4 3VGG-drop-norm  0.00595784

```

```

[ ]: ggplot(df_cnn_loss, aes(x='types', y='zero_one',group=1)) + \
      geom_point() + \
      geom_line() + \
      theme_bw(base_size=12) + ggtitle("Zero-one loss CNN") + ylab("loss")

```



```

[ ]: <ggplot: (8757990059065)>

```

6 1.5 LeNet Neural Networks

```

[ ]: input_shape = (32,32,3)
      num_classes = 10
      model = tf.keras.Sequential([
          tf.keras.layers.Conv2D(6, kernel_size=(5, 5), strides=(1, 1),
          ↪activation='tanh', input_shape=input_shape, padding="same"),
          tf.keras.layers.AveragePooling2D(pool_size=(2, 2), strides=(2, 2),
          ↪padding='valid'),
          tf.keras.layers.Conv2D(16, kernel_size=(5, 5), strides=(1, 1),
          ↪activation='tanh', padding='valid'),

```

```

        tf.keras.layers.AveragePooling2D(pool_size=(2, 2), strides=(2, 2),
        ↪padding='valid'),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(120, activation = 'tanh'),
        tf.keras.layers.Dense(84, activation = 'tanh'),
        tf.keras.layers.Dense(10, activation = 'softmax')
    ])

model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics =
    ↪['accuracy'])

model.summary()

```

Model: "sequential_16"

Layer (type)	Output Shape	Param #
conv2d_24 (Conv2D)	(None, 32, 32, 6)	456
average_pooling2d (AveragePo	(None, 16, 16, 6)	0
conv2d_25 (Conv2D)	(None, 12, 12, 16)	2416
average_pooling2d_1 (Average	(None, 6, 6, 16)	0
flatten_5 (Flatten)	(None, 576)	0
dense_32 (Dense)	(None, 120)	69240
dense_33 (Dense)	(None, 84)	10164
dense_34 (Dense)	(None, 10)	850
Total params: 83,126		
Trainable params: 83,126		
Non-trainable params: 0		

```

[ ]: history = model.fit(x_train, y_train,
                        batch_size = 32,
                        epochs = 10,
                        validation_data=(x_valid, y_valid),
                        verbose = 2

                        )

```

Epoch 1/10

WARNING:tensorflow:Method (on_train_batch_end) is slow compared to the batch update (0.219765). Check your callbacks.

1019/1019 - 8s - loss: 0.3274 - accuracy: 0.9028 - val_loss: 0.1608 - val_accuracy: 0.9496

Epoch 2/10

1019/1019 - 7s - loss: 0.0138 - accuracy: 0.9977 - val_loss: 0.0819 - val_accuracy: 0.9768

Epoch 3/10

1019/1019 - 9s - loss: 0.0014 - accuracy: 1.0000 - val_loss: 0.0839 - val_accuracy: 0.9735

Epoch 4/10

1019/1019 - 7s - loss: 6.4137e-04 - accuracy: 1.0000 - val_loss: 0.0796 - val_accuracy: 0.9779

Epoch 5/10

1019/1019 - 7s - loss: 3.2755e-04 - accuracy: 1.0000 - val_loss: 0.0749 - val_accuracy: 0.9799

Epoch 6/10

1019/1019 - 7s - loss: 1.7176e-04 - accuracy: 1.0000 - val_loss: 0.0822 - val_accuracy: 0.9790

Epoch 7/10

1019/1019 - 7s - loss: 9.9440e-05 - accuracy: 1.0000 - val_loss: 0.0775 - val_accuracy: 0.9815

Epoch 8/10

1019/1019 - 8s - loss: 5.2685e-05 - accuracy: 1.0000 - val_loss: 0.0819 - val_accuracy: 0.9813

Epoch 9/10


1019/1019 - 7s - loss: 2.9746e-05 - accuracy: 1.0000 - val_loss: 0.0737 - val_accuracy: 0.9812

Epoch 10/10

1019/1019 - 7s - loss: 1.6986e-05 - accuracy: 1.0000 - val_loss: 0.0801 - val_accuracy: 0.9822

```
[ ]: y_pred = model.predict(x_test)


# plot a random sample of test images, their predicted labels, and ground truth
fig = plt.figure(figsize=(16, 9))
for i, idx in enumerate(np.random.choice(x_test.shape[0], size=16,
    ↪replace=False)):
    ax = fig.add_subplot(4, 4, i + 1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(x_test[idx]))
    pred_idx = np.argmax(y_pred[idx])
    true_idx = np.argmax(y_test[idx])
    ax.set_title("{} ({}).format(TYPES[pred_idx], TYPES[true_idx]),
        color="green" if pred_idx == true_idx else "red"))
```

```
[ ]: plt.figure(1)

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
[ ]: y_pred = model.predict(x_test)
     zol = zero_one(y_pred, y_test)

     print("Zero-one Loss: ", zol)

     cnn_loss.append(zol)
```

Zero-one Loss: 0.016498625114573784

7 1.6 MobileNetV2

```
[ ]: model = MobileNetV2(input_shape=(32, 32, 3), alpha=1, weights=None, classes=10)
     model.compile(loss='categorical_crossentropy',
                   optimizer='adam',
                   metrics=['accuracy'])
     print(model.summary())
     print('Compiled!')
```

Model: "mobilenetv2_1.00_32"

```
-----
-----
Layer (type)                Output Shape          Param #    Connected to
=====
```

```

=====
input_12 (InputLayer)          (None, 32, 32, 3)    0
-----
Conv1_pad (ZeroPadding2D)      (None, 33, 33, 3)    0          input_12[0][0]
-----
Conv1 (Conv2D)                 (None, 16, 16, 32)   864         Conv1_pad[0][0]
-----
bn_Conv1 (BatchNormalization)  (None, 16, 16, 32)   128         Conv1[0][0]
-----
Conv1_relu (ReLU)             (None, 16, 16, 32)   0          bn_Conv1[0][0]
-----
expanded_conv_depthwise (Depthw (None, 16, 16, 32)   288
Conv1_relu[0][0]
-----
expanded_conv_depthwise_BN (Bat (None, 16, 16, 32)   128
expanded_conv_depthwise[0][0]
-----
expanded_conv_depthwise_relu (R (None, 16, 16, 32)   0
expanded_conv_depthwise_BN[0][0]
-----
expanded_conv_project (Conv2D) (None, 16, 16, 16)   512
expanded_conv_depthwise_relu[0][0]
-----
expanded_conv_project_BN (Batch (None, 16, 16, 16)   64
expanded_conv_project[0][0]
-----
block_1_expand (Conv2D)       (None, 16, 16, 96)   1536
expanded_conv_project_BN[0][0]
-----
block_1_expand_BN (BatchNormali (None, 16, 16, 96)   384
block_1_expand[0][0]
-----
block_1_expand_relu (ReLU)    (None, 16, 16, 96)   0
block_1_expand_BN[0][0]
-----

```

block_1_pad (ZeroPadding2D)	(None, 17, 17, 96)	0
block_1_expand_relu[0][0]		

block_1_depthwise (DepthwiseCon	(None, 8, 8, 96)	864
block_1_pad[0][0]		

block_1_depthwise_BN (BatchNorm	(None, 8, 8, 96)	384
block_1_depthwise[0][0]		

block_1_depthwise_relu (ReLU)	(None, 8, 8, 96)	0
block_1_depthwise_BN[0][0]		

block_1_project (Conv2D)	(None, 8, 8, 24)	2304
block_1_depthwise_relu[0][0]		

block_1_project_BN (BatchNormal	(None, 8, 8, 24)	96
block_1_project[0][0]		

block_2_expand (Conv2D)	(None, 8, 8, 144)	3456
block_1_project_BN[0][0]		

block_2_expand_BN (BatchNormali	(None, 8, 8, 144)	576
block_2_expand[0][0]		

block_2_expand_relu (ReLU)	(None, 8, 8, 144)	0
block_2_expand_BN[0][0]		

block_2_depthwise (DepthwiseCon	(None, 8, 8, 144)	1296
block_2_expand_relu[0][0]		

block_2_depthwise_BN (BatchNorm	(None, 8, 8, 144)	576
block_2_depthwise[0][0]		

block_2_depthwise_relu (ReLU)	(None, 8, 8, 144)	0
block_2_depthwise_BN[0][0]		

block_2_project (Conv2D)	(None, 8, 8, 24)	3456
block_2_depthwise_relu[0][0]		

block_2_project_BN (BatchNormal	(None, 8, 8, 24)	96
block_2_project[0][0]		

block_2_add (Add)	(None, 8, 8, 24)	0
block_1_project_BN[0][0]		
block_2_project_BN[0][0]		

block_3_expand (Conv2D)	(None, 8, 8, 144)	3456
block_2_add[0][0]		

block_3_expand_BN (BatchNormali	(None, 8, 8, 144)	576
block_3_expand[0][0]		

block_3_expand_relu (ReLU)	(None, 8, 8, 144)	0
block_3_expand_BN[0][0]		

block_3_pad (ZeroPadding2D)	(None, 9, 9, 144)	0
block_3_expand_relu[0][0]		

block_3_depthwise (DepthwiseCon	(None, 4, 4, 144)	1296
block_3_pad[0][0]		

block_3_depthwise_BN (BatchNorm	(None, 4, 4, 144)	576
block_3_depthwise[0][0]		

block_3_depthwise_relu (ReLU)	(None, 4, 4, 144)	0
block_3_depthwise_BN[0][0]		

block_3_project (Conv2D)	(None, 4, 4, 32)	4608
block_3_depthwise_relu[0][0]		

block_3_project_BN (BatchNormal	(None, 4, 4, 32)	128
block_3_project[0][0]		

```

-----
block_4_expand (Conv2D)          (None, 4, 4, 192)    6144
block_3_project_BN[0][0]

```

```

-----
block_4_expand_BN (BatchNormali (None, 4, 4, 192)    768
block_4_expand[0][0]

```

```

-----
block_4_expand_relu (ReLU)       (None, 4, 4, 192)    0
block_4_expand_BN[0][0]

```

```

-----
block_4_depthwise (DepthwiseCon (None, 4, 4, 192)    1728
block_4_expand_relu[0][0]

```

```

-----
block_4_depthwise_BN (BatchNorm (None, 4, 4, 192)    768
block_4_depthwise[0][0]

```

```

-----
block_4_depthwise_relu (ReLU)    (None, 4, 4, 192)    0
block_4_depthwise_BN[0][0]

```

```

-----
block_4_project (Conv2D)          (None, 4, 4, 32)     6144
block_4_depthwise_relu[0][0]

```

```

-----
block_4_project_BN (BatchNormal (None, 4, 4, 32)     128
block_4_project[0][0]

```

```

-----
block_4_add (Add)                 (None, 4, 4, 32)     0
block_3_project_BN[0][0]
block_4_project_BN[0][0]

```

```

-----
block_5_expand (Conv2D)          (None, 4, 4, 192)    6144
block_4_add[0][0]

```

```

-----
block_5_expand_BN (BatchNormali (None, 4, 4, 192)    768
block_5_expand[0][0]

```

```

-----
block_5_expand_relu (ReLU)       (None, 4, 4, 192)    0
block_5_expand_BN[0][0]

```

```

-----
-----
block_5_depthwise (DepthwiseCon (None, 4, 4, 192)    1728
block_5_expand_relu[0][0]
-----
-----
block_5_depthwise_BN (BatchNorm (None, 4, 4, 192)    768
block_5_depthwise[0][0]
-----
-----
block_5_depthwise_relu (ReLU)    (None, 4, 4, 192)    0
block_5_depthwise_BN[0][0]
-----
-----
block_5_project (Conv2D)          (None, 4, 4, 32)    6144
block_5_depthwise_relu[0][0]
-----
-----
block_5_project_BN (BatchNormal (None, 4, 4, 32)    128
block_5_project[0][0]
-----
-----
block_5_add (Add)                  (None, 4, 4, 32)    0
block_4_add[0][0]
block_5_project_BN[0][0]
-----
-----
block_6_expand (Conv2D)          (None, 4, 4, 192)    6144
block_5_add[0][0]
-----
-----
block_6_expand_BN (BatchNormali (None, 4, 4, 192)    768
block_6_expand[0][0]
-----
-----
block_6_expand_relu (ReLU)        (None, 4, 4, 192)    0
block_6_expand_BN[0][0]
-----
-----
block_6_pad (ZeroPadding2D)      (None, 5, 5, 192)    0
block_6_expand_relu[0][0]
-----
-----
block_6_depthwise (DepthwiseCon (None, 2, 2, 192)    1728
block_6_pad[0][0]
-----
-----
block_6_depthwise_BN (BatchNorm (None, 2, 2, 192)    768

```

block_6_depthwise[0][0]

block_6_depthwise_relu (ReLU) (None, 2, 2, 192) 0
block_6_depthwise_BN[0][0]

block_6_project (Conv2D) (None, 2, 2, 64) 12288
block_6_depthwise_relu[0][0]

block_6_project_BN (BatchNormal (None, 2, 2, 64) 256
block_6_project[0][0]

block_7_expand (Conv2D) (None, 2, 2, 384) 24576
block_6_project_BN[0][0]

block_7_expand_BN (BatchNormali (None, 2, 2, 384) 1536
block_7_expand[0][0]

block_7_expand_relu (ReLU) (None, 2, 2, 384) 0
block_7_expand_BN[0][0]

block_7_depthwise (DepthwiseCon (None, 2, 2, 384) 3456
block_7_expand_relu[0][0]

block_7_depthwise_BN (BatchNorm (None, 2, 2, 384) 1536
block_7_depthwise[0][0]

block_7_depthwise_relu (ReLU) (None, 2, 2, 384) 0
block_7_depthwise_BN[0][0]

block_7_project (Conv2D) (None, 2, 2, 64) 24576
block_7_depthwise_relu[0][0]

block_7_project_BN (BatchNormal (None, 2, 2, 64) 256
block_7_project[0][0]

block_7_add (Add) (None, 2, 2, 64) 0

block_6_project_BN[0][0]
block_7_project_BN[0][0]

block_8_expand (Conv2D) (None, 2, 2, 384) 24576
block_7_add[0][0]

block_8_expand_BN (BatchNormali (None, 2, 2, 384) 1536
block_8_expand[0][0]

block_8_expand_relu (ReLU) (None, 2, 2, 384) 0
block_8_expand_BN[0][0]

block_8_depthwise (DepthwiseCon (None, 2, 2, 384) 3456
block_8_expand_relu[0][0]

block_8_depthwise_BN (BatchNorm (None, 2, 2, 384) 1536
block_8_depthwise[0][0]

block_8_depthwise_relu (ReLU) (None, 2, 2, 384) 0
block_8_depthwise_BN[0][0]

block_8_project (Conv2D) (None, 2, 2, 64) 24576
block_8_depthwise_relu[0][0]

block_8_project_BN (BatchNormal (None, 2, 2, 64) 256
block_8_project[0][0]

block_8_add (Add) (None, 2, 2, 64) 0
block_7_add[0][0]
block_8_project_BN[0][0]

block_9_expand (Conv2D) (None, 2, 2, 384) 24576
block_8_add[0][0]

block_9_expand_BN (BatchNormali (None, 2, 2, 384) 1536
block_9_expand[0][0]

```

-----
block_9_expand_relu (ReLU)      (None, 2, 2, 384)    0
block_9_expand_BN[0][0]
-----

-----
block_9_depthwise (DepthwiseCon (None, 2, 2, 384)    3456
block_9_expand_relu[0][0]
-----

-----
block_9_depthwise_BN (BatchNorm (None, 2, 2, 384)    1536
block_9_depthwise[0][0]
-----

-----
block_9_depthwise_relu (ReLU)   (None, 2, 2, 384)    0
block_9_depthwise_BN[0][0]
-----

-----
block_9_project (Conv2D)         (None, 2, 2, 64)     24576
block_9_depthwise_relu[0][0]
-----

-----
block_9_project_BN (BatchNormal (None, 2, 2, 64)     256
block_9_project[0][0]
-----

-----
block_9_add (Add)                (None, 2, 2, 64)     0
block_8_add[0][0]
block_9_project_BN[0][0]
-----

-----
block_10_expand (Conv2D)         (None, 2, 2, 384)    24576
block_9_add[0][0]
-----

-----
block_10_expand_BN (BatchNormal (None, 2, 2, 384)    1536
block_10_expand[0][0]
-----

-----
block_10_expand_relu (ReLU)     (None, 2, 2, 384)    0
block_10_expand_BN[0][0]
-----

-----
block_10_depthwise (DepthwiseCo (None, 2, 2, 384)    3456
block_10_expand_relu[0][0]
-----

-----
block_10_depthwise_BN (BatchNor (None, 2, 2, 384)    1536
block_10_depthwise[0][0]

```

```

-----
-----
block_10_depthwise_relu (ReLU) (None, 2, 2, 384) 0
block_10_depthwise_BN[0][0]
-----
-----
block_10_project (Conv2D) (None, 2, 2, 96) 36864
block_10_depthwise_relu[0][0]
-----
-----
block_10_project_BN (BatchNorma (None, 2, 2, 96) 384
block_10_project[0][0]
-----
-----
block_11_expand (Conv2D) (None, 2, 2, 576) 55296
block_10_project_BN[0][0]
-----
-----
block_11_expand_BN (BatchNormal (None, 2, 2, 576) 2304
block_11_expand[0][0]
-----
-----
block_11_expand_relu (ReLU) (None, 2, 2, 576) 0
block_11_expand_BN[0][0]
-----
-----
block_11_depthwise (DepthwiseCo (None, 2, 2, 576) 5184
block_11_expand_relu[0][0]
-----
-----
block_11_depthwise_BN (BatchNor (None, 2, 2, 576) 2304
block_11_depthwise[0][0]
-----
-----
block_11_depthwise_relu (ReLU) (None, 2, 2, 576) 0
block_11_depthwise_BN[0][0]
-----
-----
block_11_project (Conv2D) (None, 2, 2, 96) 55296
block_11_depthwise_relu[0][0]
-----
-----
block_11_project_BN (BatchNorma (None, 2, 2, 96) 384
block_11_project[0][0]
-----
-----
block_11_add (Add) (None, 2, 2, 96) 0
block_10_project_BN[0][0]

```

block_11_project_BN[0][0]

block_12_expand (Conv2D) (None, 2, 2, 576) 55296
block_11_add[0][0]

block_12_expand_BN (BatchNormal (None, 2, 2, 576) 2304
block_12_expand[0][0]

block_12_expand_relu (ReLU) (None, 2, 2, 576) 0
block_12_expand_BN[0][0]

block_12_depthwise (DepthwiseCo (None, 2, 2, 576) 5184
block_12_expand_relu[0][0]

block_12_depthwise_BN (BatchNor (None, 2, 2, 576) 2304
block_12_depthwise[0][0]

block_12_depthwise_relu (ReLU) (None, 2, 2, 576) 0
block_12_depthwise_BN[0][0]

block_12_project (Conv2D) (None, 2, 2, 96) 55296
block_12_depthwise_relu[0][0]

block_12_project_BN (BatchNorma (None, 2, 2, 96) 384
block_12_project[0][0]

block_12_add (Add) (None, 2, 2, 96) 0
block_11_add[0][0]
block_12_project_BN[0][0]

block_13_expand (Conv2D) (None, 2, 2, 576) 55296
block_12_add[0][0]

block_13_expand_BN (BatchNormal (None, 2, 2, 576) 2304
block_13_expand[0][0]

block_13_expand_relu (ReLU)	(None, 2, 2, 576)	0
block_13_expand_BN[0][0]		

block_13_pad (ZeroPadding2D)	(None, 3, 3, 576)	0
block_13_expand_relu[0][0]		

block_13_depthwise (DepthwiseCo	(None, 1, 1, 576)	5184
block_13_pad[0][0]		

block_13_depthwise_BN (BatchNor	(None, 1, 1, 576)	2304
block_13_depthwise[0][0]		

block_13_depthwise_relu (ReLU)	(None, 1, 1, 576)	0
block_13_depthwise_BN[0][0]		

block_13_project (Conv2D)	(None, 1, 1, 160)	92160
block_13_depthwise_relu[0][0]		

block_13_project_BN (BatchNorma	(None, 1, 1, 160)	640
block_13_project[0][0]		

block_14_expand (Conv2D)	(None, 1, 1, 960)	153600
block_13_project_BN[0][0]		

block_14_expand_BN (BatchNormal	(None, 1, 1, 960)	3840
block_14_expand[0][0]		

block_14_expand_relu (ReLU)	(None, 1, 1, 960)	0
block_14_expand_BN[0][0]		

block_14_depthwise (DepthwiseCo	(None, 1, 1, 960)	8640
block_14_expand_relu[0][0]		

block_14_depthwise_BN (BatchNor	(None, 1, 1, 960)	3840
block_14_depthwise[0][0]		

block_14_depthwise_relu (ReLU)	(None, 1, 1, 960)	0
block_14_depthwise_BN[0][0]		

block_14_project (Conv2D)	(None, 1, 1, 160)	153600
block_14_depthwise_relu[0][0]		

block_14_project_BN (BatchNorma	(None, 1, 1, 160)	640
block_14_project[0][0]		

block_14_add (Add)	(None, 1, 1, 160)	0
block_13_project_BN[0][0]		
block_14_project_BN[0][0]		

block_15_expand (Conv2D)	(None, 1, 1, 960)	153600
block_14_add[0][0]		

block_15_expand_BN (BatchNormal	(None, 1, 1, 960)	3840
block_15_expand[0][0]		

block_15_expand_relu (ReLU)	(None, 1, 1, 960)	0
block_15_expand_BN[0][0]		

block_15_depthwise (DepthwiseCo	(None, 1, 1, 960)	8640
block_15_expand_relu[0][0]		

block_15_depthwise_BN (BatchNor	(None, 1, 1, 960)	3840
block_15_depthwise[0][0]		

block_15_depthwise_relu (ReLU)	(None, 1, 1, 960)	0
block_15_depthwise_BN[0][0]		

block_15_project (Conv2D)	(None, 1, 1, 160)	153600
block_15_depthwise_relu[0][0]		

block_15_project_BN (BatchNorma	(None, 1, 1, 160)	640
block_15_project[0][0]		

block_15_add (Add)	(None, 1, 1, 160)	0	
block_14_add[0][0]			
block_15_project_BN[0][0]			
block_16_expand (Conv2D)	(None, 1, 1, 960)	153600	
block_15_add[0][0]			
block_16_expand_BN (BatchNormal	(None, 1, 1, 960)	3840	
block_16_expand[0][0]			
block_16_expand_relu (ReLU)	(None, 1, 1, 960)	0	
block_16_expand_BN[0][0]			
block_16_depthwise (DepthwiseCo	(None, 1, 1, 960)	8640	
block_16_expand_relu[0][0]			
block_16_depthwise_BN (BatchNor	(None, 1, 1, 960)	3840	
block_16_depthwise[0][0]			
block_16_depthwise_relu (ReLU)	(None, 1, 1, 960)	0	
block_16_depthwise_BN[0][0]			
block_16_project (Conv2D)	(None, 1, 1, 320)	307200	
block_16_depthwise_relu[0][0]			
block_16_project_BN (BatchNorma	(None, 1, 1, 320)	1280	
block_16_project[0][0]			
Conv_1 (Conv2D)	(None, 1, 1, 1280)	409600	
block_16_project_BN[0][0]			
Conv_1_bn (BatchNormalization)	(None, 1, 1, 1280)	5120	Conv_1[0][0]
out_relu (ReLU)	(None, 1, 1, 1280)	0	Conv_1_bn[0][0]

```

global_average_pooling2d_1 (Glo (None, 1280)          0          out_relu[0][0]
-----
Logits (Dense)                (None, 10)          12810
global_average_pooling2d_1[0][0]
=====
Total params: 2,270,794
Trainable params: 2,236,682
Non-trainable params: 34,112
-----
None
Compiled!

```

```
[ ]: history = model.fit(x_train,y_train,
    batch_size = 32,
    epochs=10,
    validation_data=(x_valid, y_valid),
    verbose=2 )
```

```

Train on 32607 samples, validate on 8724 samples
Epoch 1/10
- 152s - loss: 0.6746 - accuracy: 0.7860 - val_loss: 2.2316 - val_accuracy:
0.1977
Epoch 2/10
- 148s - loss: 0.1768 - accuracy: 0.9496 - val_loss: 2.2468 - val_accuracy:
0.1977
Epoch 3/10
- 148s - loss: 0.1177 - accuracy: 0.9681 - val_loss: 2.2137 - val_accuracy:
0.1977
Epoch 4/10
- 147s - loss: 0.1166 - accuracy: 0.9705 - val_loss: 2.1745 - val_accuracy:
0.1977
Epoch 5/10
- 147s - loss: 0.0993 - accuracy: 0.9756 - val_loss: 0.8243 - val_accuracy:
0.7942
Epoch 6/10
- 149s - loss: 0.0763 - accuracy: 0.9815 - val_loss: 0.9804 - val_accuracy:
0.7658
Epoch 7/10
- 146s - loss: 0.0685 - accuracy: 0.9836 - val_loss: 2.2230 - val_accuracy:
0.6530
Epoch 8/10
- 149s - loss: 0.0582 - accuracy: 0.9859 - val_loss: 9.0434 - val_accuracy:
0.4411
Epoch 9/10
- 148s - loss: 0.0601 - accuracy: 0.9856 - val_loss: 4.6619 - val_accuracy:

```



0.6551

Epoch 10/10

- 149s - loss: 0.0429 - accuracy: 0.9905 - val_loss: 2.4783 - val_accuracy:
0.7839

```
[ ]: y_pred = model.predict(x_test)

# plot a random sample of test images, their predicted labels, and ground truth
fig = plt.figure(figsize=(16, 9))
for i, idx in enumerate(np.random.choice(x_test.shape[0], size=16,
↪replace=False)):
    ax = fig.add_subplot(4, 4, i + 1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(x_test[idx]))
    pred_idx = np.argmax(y_pred[idx])
    true_idx = np.argmax(y_test[idx])
    ax.set_title("{} ({}).format(TYPES[pred_idx], TYPES[true_idx]),
                color="green" if pred_idx == true_idx else "red"))
```




```
[ ]: plt.figure(1)

plt.subplot(211)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
```

```
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
[ ]: y_pred = model.predict(x_test)
      zol = zero_one(y_pred, y_test)
      cnn_loss.append(zol)
      print("Zero-one Loss: ", zol)
```

Zero-one Loss: 0.18973418881759854

8 1.7 Summary results

```
[ ]: df_cnn_loss = pd.DataFrame(  
    columns = ['type', 'zero_one']  
)  
  
types = ["1VGG", "2VGG", "3VGG", "3VGG-drop",  
    ↪ "3VGG-drop-norm", "LeNet", "MobileNetV2"]  
df_cnn_loss = pd.DataFrame( data = [types, cnn_loss], index = ['types',  
    ↪ 'zero_one']).T  
  
df_new = df_loss.rename(columns={'components': 'types'})  
df_new['types'] = df_new['types'].astype(str)  
  
[ ]: frames = [df_new, df_cnn_loss]  
result = pd.concat(frames)  
result = result.sort_values(by=['zero_one'])  
  
result = result.reset_index(drop=True)  
  
[ ]: plt.barh( result['types'].values, result['zero_one'].values , align='center',  
    ↪ alpha=0.5)  
plt.xlabel('Loss')  
plt.ylabel("Loss")  
plt.title('Model Loss summary')  
  
plt.show()
```

